

Distributed Systems Assignment

s1140740

1 Prove $a \rightarrow b \iff V(a) \leq V(b)$.

1.1 Proof of $a \rightarrow b \implies V(a) \leq V(b)$ by contradiction

Assume $a \rightarrow b$ and $V(a) > V(b)$. This means that $\exists j. V(a)[j] > V(b)[j]$. For this to be true there had to be two successive events $t-1$ and t between a and b such that $V(t-1)[j] > V(t)[j]$.

There are two cases:

1. Events $t-1$ and t were successive events in same process. However, by definition after every local event of a process i , $V_i[i] = V_i[i] + 1$.
2. Event $t-1$ was the “send” event of message M from process i and event t was the “receive” event of message M at process j . By definition, $V(t-1)$ was attached to M . Furthermore, after process j receives the message as event t it performs the following steps:
 1. $V_j[k] = \max(V_j[k], V_i[k])$, for $k = 1, 2, \dots, n$
 2. $V_j[j] = V_j[j] + 1$

Both cases imply $V(t-1) < V(t)$ and thus we arrive at a contradiction.

1.2 Proof of $V(a) \leq V(b) \implies a \rightarrow b$ by contradiction

2 Inductive proof on the position of the request

2.1 Base case

Request is at position 1 in the queue and thus the process can access the resource and satisfy the request.

2.2 Induction hypothesis

The request eventually gets satisfied at position k .

2.3 Inductive step

Request is at position $k + 1$. Let the request at first position belong to process i and let us call the request R_i . Since R_i is at the first position in the queue, all the previous processes accessing the resource are finished with it, otherwise we would have their requests in our queue before R_i (we add a request to the queue whenever we get a **REQUEST** message and remove it only once we receive a **RELEASE** message for it). Thus, there are three cases:

1. Process i is currently accessing the resource. Since we assume processes do not fail, this means that it will eventually finish accessing it and when it does it will send us a **RELEASE** message and R_i will be removed from our queue and our request will be in position k .
2. Process i has already finished accessing the resource. This implies we have not yet received the **RELEASE** message. Channels do not fail so we will eventually receive the message and remove R_i from our queue and our request will be in position k .
3. Process i has not started accessing the resource. This implies R_i is not yet at the top of the queue of process i (otherwise it would just access the resource). However, since we have shown that no process with a request before R_i can be accessing the resource this means that process i just has not received the **RELEASE** message from the last process accessing the resource. Once it receives this message it will pop that process's request of its queue and start accessing the resource. Logic in case 1 can then be followed to show that our request will advance to position k .

Using induction hypothesis we can show the request will eventually get satisfied in all three cases.

3 Weighted diameter

The weighted diameter of this graph is 7. The path realising this diameter is $A \rightarrow C \rightarrow E \rightarrow G \rightarrow H$.

If the graph was unweighted, the diameter would be 4 and the corresponding path would be $A \rightarrow C \rightarrow F \rightarrow G \rightarrow I$.

4 Prim's algorithm

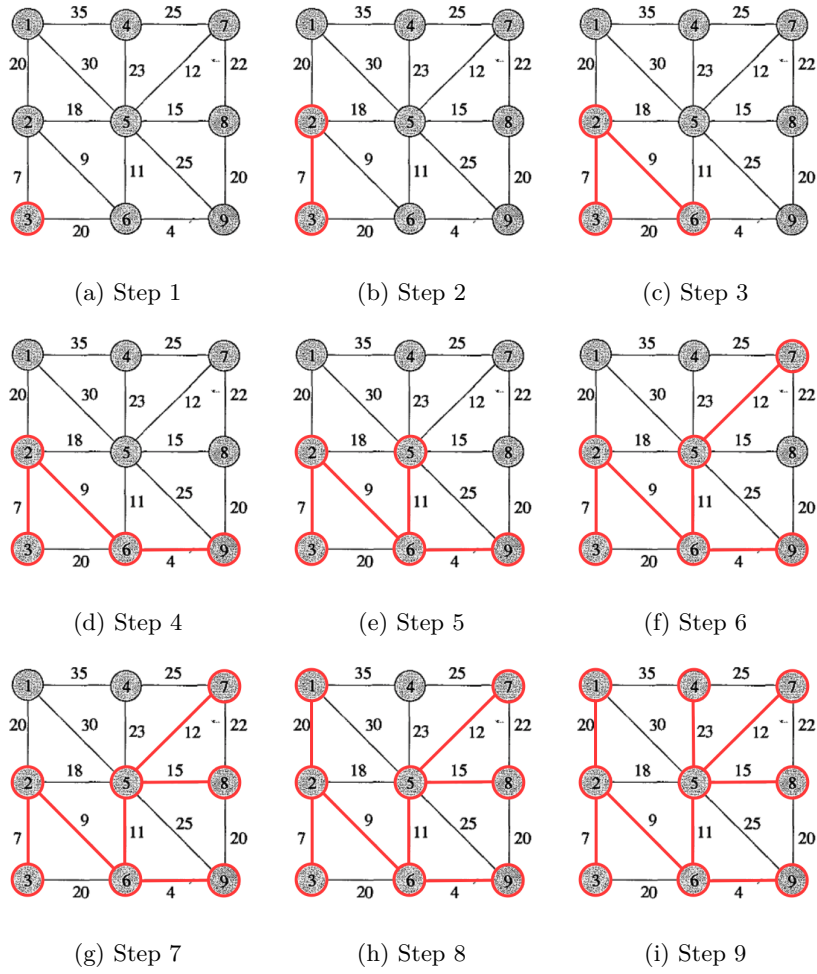


Figure 1: Prim's algorithm